



HW2 Camera Relocalization

Due: 2023/10/24 11:59 AM

3DCV 2023

Email: 3dcv@csie.ntu.edu.tw


GitHub Classroom: <https://classroom.github.com/a/cKfvp3Eo>

GitHub Registration: <https://forms.gle/R9JBIAAehcYyvoUu9>

Outline

The goal of this homework is to realize how a camera re-localization system works.

- Introduction
- Dataset
- Problem 1: 2D-3D Matching (Q1-1 ~ Q1-3)
- Problem 2: Augmented Reality (Q2-1 ~ Q2-2)
- Bonus List
- Grading Policy



Introduction

- **Camera Relocalization:** Determine the camera pose from the visual scene representation. In other words, the scene is **seen** (and modeled) **beforehand**. Now, given a query image that is taken in this environment, we are able to find out where this image is taken.




Query Image


Training Data (Database Image)



Projection



Recover




Real World (3D Information)



Introduction

- **One-shot relocation:** focus on finding the pose of still image.
- **Temporal camera relocation:** estimates the poses of every frame in the video sequence




Methodology

- Common strategies for camera relocalization. Note that there are some approaches utilize hybrid models to increase the efficiency and robustness.
- Metric localization can only be achieved by machine (or deep) learning models.

2D-3D Matching Localization

3D scene models with local
feature descriptors



Metric Localization (Absolute Pose Regression)

regresses the position and
orientation of the camera





Image Retrieval Localization

pose approximation from top-k
similar images



Welcome to the NTU Front Gate

- We collect multiple images of the NTU front gate, and reconstruct its 3D point cloud model via structure from motion.

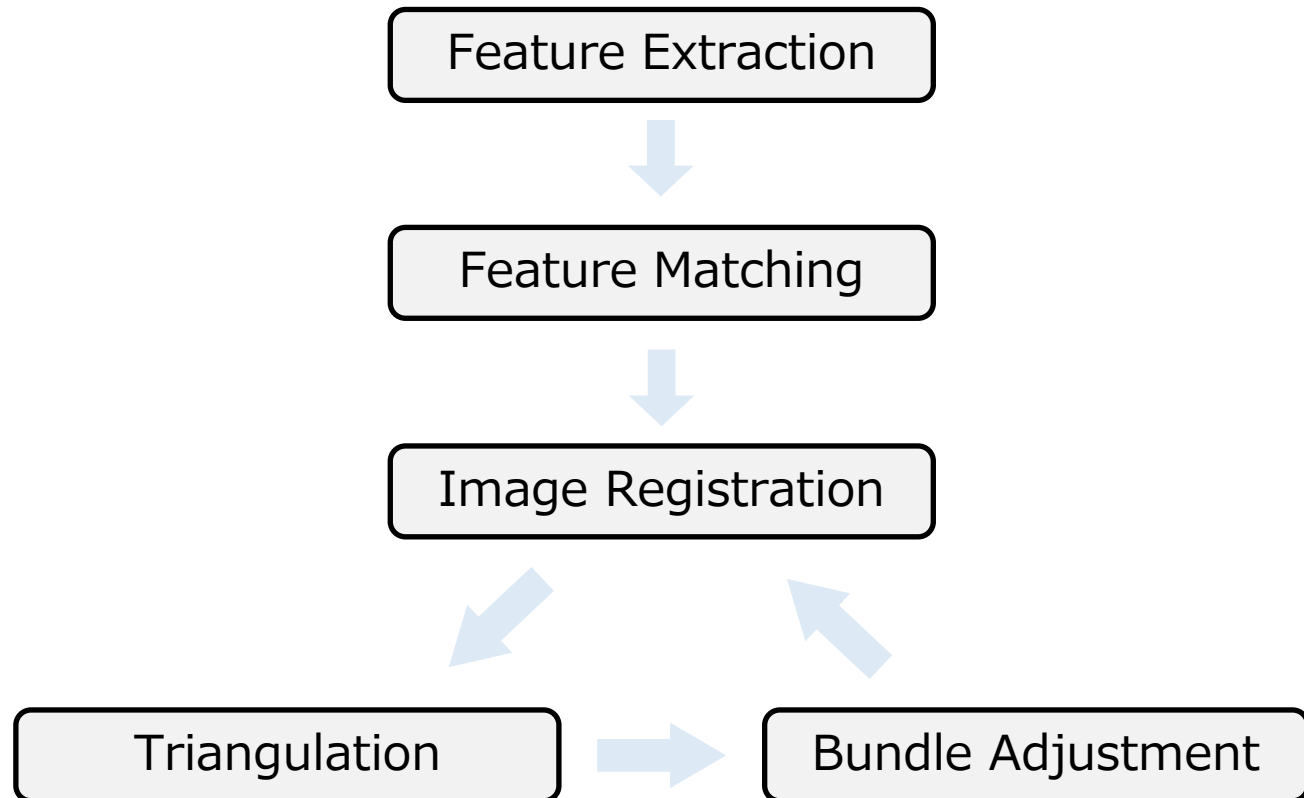


About Dataset

- 293 color images (1920x1080x3): 163 images for training, 130 images for testing
- 111,518 points (in world coordinate) with 682,467 local image descriptors




Dataset images



About Dataset

- 293 color images (1920x1080x3): 163 images for training, 130 images for testing
- 111,518 points (in world coordinate) with 682,467 local image descriptors



Data/image.pkl

⚠ The pose of an image is represented as the projection **from world to the camera** coordinate system. That is, $p = K[R | T]X$.

			Camera Position(x,y,z)			Rotation (in quaternion)			
IMAGE_ID		NAME	TX	TY	TZ	QW	QX	QY	QZ
0	1	train_img100.jpg	-3.12923	-0.273371	3.17218	0.969363	-0.003488	0.244797	0.019927
1	2	train_img104.jpg	-3.10598	-0.264036	3.12049	0.972423	-0.005048	0.232322	0.019880
2	3	train_img108.jpg	-3.06986	-0.270274	3.08285	0.975032	-0.004203	0.221007	0.021220
3	4	train_img112.jpg	-3.02027	-0.290710	3.07195	0.976940	-0.003627	0.212336	0.022091
4	5	train_img116.jpg	-2.98028	-0.307973	3.05439	0.979017	-0.002989	0.202524	0.022389
...
288	289	valid_img75.jpg	-2.86676	-0.366566	3.79563	0.931094	0.002295	0.363172	0.034118
289	290	valid_img80.jpg	-2.86618	-0.323873	3.72239	0.937160	-0.001973	0.347476	0.031431
290	291	valid_img85.jpg	-2.91426	-0.300918	3.59808	0.945271	-0.004261	0.325035	0.028210
291	292	valid_img90.jpg	-2.99320	-0.267023	3.46717	0.954254	-0.004443	0.298019	0.023733
292	293	valid_img95.jpg	-3.08001	-0.259334	3.30072	0.962891	-0.003862	0.269045	0.021006

293 rows × 9 columns

⚠ Note that the order is (QW, QX, QY, QZ)

Data/point_desc.pkl

- point_desc.pkl

Source Info

128D Descriptors

	POINT_ID	IMAGE_ID	XY	DESCRIPTORS
0	1	1	[94.94650268554688, 284.02899169921875]	[46, 43, 12, 11, 10, 5, 19, 37, 24, 16, 8, 9, ...]
1	1	2	[99.05780029296875, 290.6889953613281]	[39, 42, 34, 14, 15, 12, 13, 31, 29, 11, 8, 7, ...]
2	1	3	[110.51899719238281, 291.7560119628906]	[47, 57, 39, 12, 12, 11, 9, 20, 43, 26, 13, 7, ...]
3	1	4	[131.70199584960938, 286.4880065917969]	[38, 58, 39, 12, 11, 11, 13, 16, 35, 20, 12, 8, ...]
4	1	7	[156.52499389648438, 279.2149963378906]	[32, 38, 31, 19, 15, 6, 11, 32, 28, 14, 6, 10, ...]
...
1234453	129081	276	[816.5590209960938, 353.6910095214844]	[28, 20, 11, 16, 23, 18, 22, 25, 42, 11, 8, 24, ...]
1234454	129081	278	[892.0490112304688, 384.6050109863281]	[30, 30, 15, 22, 28, 14, 15, 23, 47, 13, 10, 2, ...]
1234455	129081	279	[965.5770263671875, 397.2950134277344]	[29, 22, 12, 18, 28, 16, 20, 30, 40, 12, 9, 27, ...]
1234456	129081	280	[1039.56005859375, 405.864990234375]	[27, 24, 14, 15, 26, 16, 25, 33, 45, 12, 10, 2, ...]
1234457	129081	280	[1045.989990234375, 404.6090087890625]	[23, 38, 24, 33, 28, 7, 3, 7, 52, 12, 12, 26, ...]

⚠ If Point_ID is -1, then its 3D position is not available.

Data/train.pkl

- train.pkl 3D Point Position(x,y,z)

Source Info

128D Descriptors

POINT_ID		XYZ	RGB	IMAGE_ID		XY	DESCRIPTORS
0	1	[1.6093346, -1.1848674, 1.610395]	[87, 87, 77]	1		[94.94650268554688, 284.02899169921875]	[46, 43, 12, 11, 10, 5, 19, 37, 24, 16, 8, 9, ...]
1	1	[1.6093346, -1.1848674, 1.610395]	[87, 87, 77]	2		[99.05780029296875, 290.6889953613281]	[39, 42, 34, 14, 15, 12, 13, 31, 29, 11, 8, 7, ...]
2	1	[1.6093346, -1.1848674, 1.610395]	[87, 87, 77]	3		[110.51899719238281, 291.7560119628906]	[47, 57, 39, 12, 12, 11, 9, 20, 43, 26, 13, 7, ...]
3	1	[1.6093346, -1.1848674, 1.610395]	[87, 87, 77]	4		[131.70199584960938, 286.4880065917969]	[38, 58, 39, 12, 11, 11, 13, 16, 35, 20, 12, 8, ...]
4	1	[1.6093346, -1.1848674, 1.610395]	[87, 87, 77]	7		[156.52499389648438, 279.2149963378906]	[32, 38, 31, 19, 15, 6, 11, 32, 28, 14, 6, 10, ...]
...
682463	129081	[0.66382873, -1.3121917, 5.433149]	[33, 30, 28]	141		[834.9459838867188, 363.7510070800781]	[32, 26, 15, 19, 28, 14, 18, 30, 37, 12, 11, 2, ...]
682464	129081	[0.66382873, -1.3121917, 5.433149]	[33, 30, 28]	142		[867.6019897460938, 366.8039855957031]	[33, 16, 6, 11, 25, 16, 18, 36, 41, 10, 7, 23, ...]
682465	129081	[0.66382873, -1.3121917, 5.433149]	[33, 30, 28]	144		[981.5599975585938, 398.8039855957031]	[25, 14, 7, 12, 27, 21, 24, 28, 50, 13, 8, 24, ...]
682466	129081	[0.66382873, -1.3121917, 5.433149]	[33, 30, 28]	145		[1039.56005859375, 405.864990234375]	[27, 24, 14, 15, 26, 16, 25, 33, 45, 12, 10, 2, ...]
682467	129081	[0.66382873, -1.3121917, 5.433149]	[33, 30, 28]	145		[1045.989990234375, 404.6090087890625]	[23, 38, 24, 33, 28, 7, 3, 7, 52, 12, 12, 26, ...]

682468 rows × 6 columns

About Dataset: Camera Parameters

- Review the Pinhole camera model:

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \approx \begin{bmatrix} f_x & s & o_x \\ 0 & f_y & o_y \\ 0 & 0 & 1 \end{bmatrix} [R \quad t] \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

- Intrinsic Parameters:

$$K = \begin{bmatrix} f_x & s & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 1868.27 & 0 & 540 \\ 0 & 1869.18 & 960 \\ 0 & 0 & 1 \end{bmatrix}$$

- Distortion Parameters (Brown-Conrady Model):

$$D = [k_1 \quad k_2 \quad p_1 \quad p_2] = [0.0847023, -0.192929, -0.000201144, -0.000725352]$$

Problem1



Q1-1 For each validation image, compute its camera pose with respect to world coordinate. Find the 2D-3D correspondence by descriptor matching, and solve the camera pose. Implement at least one kind of algorithm that solves a PnP problem. Briefly explain your implementation and write down the pseudo code in your report.

Notes:

- Expected Solution: P3P + RANSAC. You have to implement RANSAC by yourself.
- You cannot use calib3d module in OpenCV. That is, solvePnP and solvePnPRansac is forbidden. However, you are encouraged to try them beforehand.
- You may also try DLT, EPnP, AP3P, or any kinds of solutions.

Problem1

Q1-2 For each camera pose you calculated, compute the median pose error (translation, rotation) with respect to ground truth camera pose. Provide some discussion.


Notes:

- Translation: median of all absolute pose differences (Euclidean Distance).

$$t_e = \|\mathbf{t} - \hat{\mathbf{t}}\|_2$$

- Rotation: median of relative rotation angle between estimation and ground-truth.
(1. Find out the relative rotation and represent it as axis angle representation.
2. Report the median of angles.)

$$\mathcal{R} = R_e \hat{\mathcal{R}}$$




Problem1

Q1-3 For each camera pose you calculated, plot the trajectory and camera poses along with 3d point cloud model using Open3D. Explain how you draw and provide some discussion.

Notes:

- Draw the camera pose as a quadrangular pyramid, where the apex is the position of the optical center, and the normal of base is the bearing (orientation) of the camera.





Problem2

Q2-1 With camera intrinsic and extrinsic parameters, place a virtual cube in the the validation image sequences to create an Augmented Reality video. Draw the virtual cube as a point set with different colors on its surface. Implement a simply but efficient painter's algorithm to determine the order of drawing.

Notes:

- You don't have to consider whether virtual cube will be occluded.
- Manually select the location, orientation, and scale of the virtual cube. (We provide a code that allows you to adjust the cube by keyboard.)
- Painter's Algorithm:
 1. Sort each voxel by depth
 2. Place each voxel from the furthest to the closest



Sample Code

- You should read the pickle files with **pandas**.

```
>>> import pandas as pd
>>> images_df = pd.read_pickle("dataframes/images.pkl")
```

- You may use **Scipy** to deal with 3D rotation representations.

```
>>> from scipy.spatial.transform import Rotation as R
>>> r = R.from_quat([0, 0, np.sin(np.pi/4), np.cos(np.pi/4)])
>>> r.as_rotvec()
array([0. , 0. , 1.57079633])
```

Parameters: `quat` : *array_like, shape (N, 4) or (4,)*

Each row is a (possibly non-unit norm) quaternion in **scalar-last (x, y, z, w) format**. Each quaternion will be normalized to unit norm.

⚠ Be aware of the order.

Returns: `rotation` : *Rotation instance*

Object containing the rotations represented by input quaternions.

Introduction to Open3D

- Install open3D `pip install open3d`
- Basic manipulation in open3D (Example Drawing):

```
points = [[0, 0, 0], [1, 0, 0], [0, 1, 0], [1, 1, 0],  
          [0, 0, 1], [1, 0, 1], [0, 1, 1], [1, 1, 1]]  
lines = [[0, 1], [0, 2], [1, 3], [2, 3], [4, 5], [4, 6],  
         [5, 7], [6, 7], [0, 4], [1, 5], [2, 6], [3, 7]]
```

```
import open3d as o3d  
line_set = o3d.geometry.LineSet()  
line_set.points = o3d.utility.Vector3dVector(points)  
line_set.lines = o3d.utility.Vector2iVector(lines)
```

```
vis = o3d.visualization.Visualizer()  
vis.create_window()  
vis.add_geometry(line_set) o3d.visualization.ViewControl.set_zoom(vis.get_view_control(), 0.8)  
vis.run()
```

⚠ Please refer to the document to find the property you need.

Bonus List



To get extra credits, you can try the following things:
(including, but not limited to)

- **Local Features:** Try different kinds of local features (including deep features)
- **Make it faster:** Come up with faster matching or image registration strategy. (prioritized matching, approximate nearest neighbor, coarse-to-fine strategy, image retrieval, ...)
- **Make it more accurate:** Make the pose estimation more accurate. (Different PnP solving methods, outlier rejection strategies, ...)
- **Absolute Pose Regression:** Train a deep neural network to regress the absolute camera pose. (For example, PoseNet: A Convolutional Network for Real-Time 6-DOF Camera Relocalization, ICCV 2015)

Grading



- We will evaluate both **the functionality of the code** and **the quality of the report**.
- **[*] Youtube link :**
 - You should record your demonstration, including the start time and the GitHub clone action
 - Example : <https://youtu.be/-VnjVda7c8o?si=zowfe7vjvCMMFrOk>
- **Functionality:** Can it run? How's the performance?
- **Quality:** theoretical/experimental analysis, observation, discussion, ...
- Note that it **might be curved** based on overall performance of students.
- Grade
 - Meet the basic requirement (programming & report) → A
 - Basic requirement + advanced studies (programming & report) → A+

Grading Policies



- Push your code and report to the GitHub classroom.
- Programming Languages: Python (Python \geq 3.8), (C++)
- Report Format: PDF or Markdown
(Warning for Markdown users: Latex equations cannot be rendered properly in GitHub)
- Late Submission: **-10% from your score** / day
- Plagiarism: You have to **write your own codes**.
- Discussion: We encourage you to discuss with your classmates, but remember to **mention their names and contributions in the report**.

Thanks

If you have any question, please email 3dcv@csie.ntu.edu.tw